

CAP6610 Machine Learning, Spring 2022

Midterm 1

2/22/2022

1. (15 points) *Fitting a piecewise linear function to data.* We are given some samples $(x_i, y_i), i = 1, \dots, n$, with x_i, y_i scalars, with a function of the form

$$\hat{y} = f(x) = \theta_1 x + \theta_2 \max(0, x) + \theta_3 \min(0, x),$$

where $\boldsymbol{\theta} = (\theta_1, \theta_2, \theta_3)$ contains the parameters. The function is affine for $x < 0$ and also for $x > 0$, and it is continuous at $x = 0$ with value θ_1 . Such a function is called *piecewise affine*, or more commonly *piecewise linear*, with a single knot or kink point at 0. We choose the parameters using the least squares regression on the given data points.

Consider the specific data set of (x, y) pairs

$$(-2, 3), (-1, 1), (0, 1), (1, 3), (2, 2)$$

Give the matrix $\boldsymbol{\Phi}$ and vector $\boldsymbol{\psi}$ for which the sum of the squares of the fitting errors is equal to $\|\boldsymbol{\Phi}\boldsymbol{\theta} - \boldsymbol{\psi}\|^2$. We want the explicit numerical values of $\boldsymbol{\Phi}$ and $\boldsymbol{\psi}$. Then use any programming language to solve this least squares problem and give the numerical values of the optimal $\boldsymbol{\theta}$.

2. (15 points) Consider the document classification problem. The training data is a set of documents with their correct labels. Suppose the docs are categorized into k classes, and y_i denotes the class that doc i belongs to. Each doc is transformed into the “bag-of-words” representation, say \mathbf{x}_i for the i th doc, meaning the j th entry of \mathbf{x}_i represents the number of times term j appears in doc i . We want to estimate the probability $p(y|\mathbf{x})$ in order to design a document classifier. We use a generative approach: according to Bayes’ rule, $p(y|\mathbf{x})$ is proportional to $p(y)p(\mathbf{x}|y)$:

- (a) Describe how to estimate $p(y)$
 - (b) Assume $p(\mathbf{x}|y)$ follows a multinomial distribution; for each class c , the parameter for the multinomial distribution is a nonnegative vector \mathbf{p}_c that sums to one. Describe how to estimate \mathbf{p}_c .
 - (c) Write the resulting classifier in the form $\hat{y} = \max_c(\mathbf{w}_c^\top \mathbf{x} + \beta_c)$. Explain how to obtain \mathbf{w}_c and β_c from $p(y)$ and $p(\mathbf{x}|y)$.
3. (20 points) *Some sets of probability distributions.* Let x be a real-valued random variable with sample space $\{a_1, \dots, a_k\}$ where $a_1 \leq a_2 \leq \dots \leq a_k$. This can be viewed as a categorical random variable with each category assigned a real value. Let $\Pr[x = a_i] = p_i$, then the vector \mathbf{p} satisfies $\mathbf{p} \geq 0$ and $\mathbf{1}^\top \mathbf{p} = 1$, i.e., it lies in the probability simplex Δ . Which of the following conditions are convex in \mathbf{p} ? (That is, defines a set of \mathbf{p} that is a convex set.) Circle the correct answers. You do not need to justify your responses.
- (a) $\alpha \leq \mathbb{E}[f(x)] \leq \beta$, where $\mathbb{E}[f(x)]$ is the expected value of $f(x)$, i.e., $\mathbb{E}[f(x)] = \sum_{i=1}^k p_i f(a_i)$. (The function $f: \mathbb{R} \rightarrow \mathbb{R}$ is given.)
 - (b) $\Pr[x > \alpha] \leq \beta$.

(c) $E[x^3] \leq \alpha E[x]$.

(d) $\text{var}(x) \leq \alpha$.

(e) $\text{var}(x) \geq \alpha$.

4. (17 points) *Maximum likelihood estimation for exponential family.* A probability distribution or density on a set \mathcal{X} , parameterized by $\boldsymbol{\theta} \in \mathbb{R}^m$, is called an *exponential family* if it has the form

$$p(\mathbf{x}; \boldsymbol{\theta}) = a(\boldsymbol{\theta}) \exp(\boldsymbol{\theta}^\top \boldsymbol{\phi}(\mathbf{x})),$$

for $\mathbf{x} \in \mathcal{X}$, where $\boldsymbol{\phi} : \mathcal{X} \rightarrow \mathbb{R}^m$, and $a(\boldsymbol{\theta})$ is a normalization function. Here we interpret as a density function when \mathcal{X} is a continuous set, and a probability distribution if \mathcal{X} is discrete. Thus we have

$$a(\boldsymbol{\theta}) = \left(\int_{\mathcal{X}} \exp(\boldsymbol{\theta}^\top \boldsymbol{\phi}(\mathbf{x})) d\mathbf{x} \right)^{-1}$$

when $p(\mathbf{x}; \boldsymbol{\theta})$ is a density, and

$$a(\boldsymbol{\theta}) = \left(\sum_{\mathbf{x} \in \mathcal{X}} \exp(\boldsymbol{\theta}^\top \boldsymbol{\phi}(\mathbf{x})) \right)^{-1}$$

when $p(\mathbf{x}; \boldsymbol{\theta})$ represents a distribution. We consider only values of $\boldsymbol{\theta}$ for which the integral or sum above is finite. Many families of distributions have this form, for appropriate choice of the parameter $\boldsymbol{\theta}$ and function $\boldsymbol{\phi}$.

Show that for any $\mathbf{x} \in \mathcal{X}$, the log-likelihood function $\log p(\mathbf{x}; \boldsymbol{\theta})$ is concave in $\boldsymbol{\theta}$. This means that maximum-likelihood estimation for an exponential family leads to a convex optimization problem. You don't have to give a formal proof of concavity of $\log p(\mathbf{x}; \boldsymbol{\theta})$ in the general case: You can just consider the case when \mathcal{X} is finite, and state that the other cases (discrete but infinite \mathcal{X} , continuous \mathcal{X}) can be handled by taking limits of finite sums.

5. (17 points) *Fitting with censored data.* In some experiments there are two kinds of measurements or data available: The usual ones, in which you get a number (say), and censored data, in which you don't get the specific number, but are told something about it, such as a lower bound. A classic example is a study of lifetimes of a set of subjects (say, laboratory mice). For those who have died by the end of data collection, we get the lifetime. For those who have not died by the end of data collection, we do not have the lifetime, but we do have a lower bound, i.e., the length of the study. These are the censored data values.

We wish to fit a set of data points $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$, with $\mathbf{x}_i \in \mathbb{R}^m$ and $y \in \mathbb{R}$, using a linear model of the form $y \approx \mathbf{x}^\top \mathbf{w}$. The vector $\mathbf{w} \in \mathbb{R}^m$ is the model parameter, which we wish to choose. We will use a least-squares criterion, i.e., choose \mathbf{w} to minimize

$$\sum_{i=1}^n (y_i - \mathbf{x}_i^\top \mathbf{w})^2.$$

Here is the tricky part: some of the values of y_i are censored; for these entries, we have only a (given) lower bound. We will re-order the data so that y_1, \dots, y_p are given (i.e., uncensored), while y_{p+1}, \dots, y_n are all censored, i.e., unknown, but larger than D , a given number. All the values of \mathbf{x}_i are known.

Formulate an optimization problem to find \mathbf{w} (the model parameter) and y_{p+1}, \dots, y_n (the censored data values). Be sure to clarify the three essential components of an optimization problem: the loss function to be minimized, the optimization variables, and constraints (if any). Is it a convex problem?

6. (16 points) *Poisson regression*. Consider a regression problem with training data $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)$. Again our predictor takes the form $\hat{y} = \boldsymbol{\phi}^\top \boldsymbol{\theta}$, where $\boldsymbol{\phi}$ is a feature vector obtained from \mathbf{x} , and the coefficients $\boldsymbol{\theta}$ is learned from a supervised learning model.

In many cases the target output y is always a positive integer, so it makes sense to assume that it follows a Poisson distribution. Specifically, we assume that the conditional probability $p(y|\mathbf{x})$ follows a Poisson distribution with parameter $\lambda = e^{\boldsymbol{\phi}^\top \boldsymbol{\theta}}$. Recall that the Poisson pmf with parameter λ is

$$p(y) = \frac{\lambda^y e^{-\lambda}}{y!}.$$

Formulate an optimization problem to learn $\boldsymbol{\theta}$. Is it a convex optimization problem?

Hint. We will maximize the log-likelihood of the training data, assuming the samples are independent.